

HW8

Irene Yang

4/18/2018

K-nearest neighbor

Let's try a variation on the NHANES data set again.

```
library(tidyverse)

## — Attaching packages — tidyverse 1.2.1 —

## [ggplot2 2.2.1] [purrr 0.2.4]
## [tibble 1.4.2] [dplyr 0.7.4]
## [tidyr 0.8.0] [stringr 1.3.0]
## [readr 1.1.1] [forcats 0.3.0]

## — Conflicts —
tidyverse_conflicts() —
## [dplyr::filter()] masks stats::filter()
## [dplyr::lag()] masks stats::lag()

library(class)
library(rpart)
library(NHANES)
library(RColorBrewer)
library(plot3D)
library(parallel)
library(randomForestSRC)

##
## randomForestSRC 2.5.1
##
## Type rfsrc.news() to see new features, changes, and bug fixes.
##

##
## Attaching package: 'randomForestSRC'

## The following object is masked from 'package:purrr':
##
## partial

library(ggRandomForests)

##
## Attaching package: 'ggRandomForests'
```

```
## The following object is masked from 'package:randomForestSRC':
##
##     partial.rfsrc

library(mosaic)

## Loading required package: lattice

## Loading required package: ggformula

##
## New to ggformula? Try the tutorials:
##   learnr::run_tutorial("introduction", package = "ggformula")
##   learnr::run_tutorial("refining", package = "ggformula")

## Loading required package: mosaicData

## Loading required package: Matrix

##
## Attaching package: 'Matrix'

## The following object is masked from 'package:tidyr':
##
##     expand

##
## The 'mosaic' package masks several functions from core packages in
## order to add
## additional features. The original behavior of these functions
## should not be affected by this.
##
## Note: If you use the Matrix package, be sure to load it BEFORE
## loading mosaic.

##
## Attaching package: 'mosaic'

## The following object is masked from 'package:Matrix':
##
##     mean

## The following objects are masked from 'package:dplyr':
##
##     count, do, tally

## The following object is masked from 'package:purrr':
##
##     cross

## The following objects are masked from 'package:stats':
##
```

```
##      binom.test, cor, cor.test, cov, fivenum, IQR, median,
##      prop.test, quantile, sd, t.test, var
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      max, mean, min, prod, range, sample, sum
```

```
# Create the NHANES dataset again
```

Create the NHANES dataset again, just like we did in class, only using sleep trouble (variable name = SleepTrouble) as the dependent variable, instead of SleepTrouble.

```
# Create the NHANES dataset again
```

```
people2 <- NHANES %>% dplyr::select(Age, Gender, SleepTrouble, BMI,
HHIncome, PhysActive)
#%>% na.omit()
```

```
glimpse(people2)
```

```
## Observations: 10,000
```

```
## Variables: 6
```

```
## $ Age          <int> 34, 34, 34, 4, 49, 9, 8, 45, 45, 45, 66, 58,
54, ...
```

```
## $ Gender       <fct> male, male, male, male, female, male, male,
femal...
```

```
## $ SleepTrouble <fct> Yes, Yes, Yes, NA, Yes, NA, NA, No, No, No, No,
N...
```

```
## $ BMI          <dbl> 32.22, 32.22, 32.22, 15.30, 30.57, 16.82,
20.64, ...
```

```
## $ HHIncome     <fct> 25000-34999, 25000-34999, 25000-34999, 20000-
2499...
```

```
## $ PhysActive   <fct> No, No, No, NA, No, NA, NA, Yes, Yes, Yes, Yes,
Y...
```

Problem 1

What is the marginal distribution of sleep trouble?

```
tally(~ SleepTrouble, data = people2, format = "percent")
```

```
## SleepTrouble
```

```
##      No      Yes <NA>
```

```
## 57.99 19.73 22.28
```

Recall from our prior work, the packages work better if the dataset is a dataframe, and the variables are numeric.

```
class(people2)
```

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

```

# Convert back to dataframe
people2 <- as.data.frame(people2)
glimpse(people2)

## Observations: 10,000
## Variables: 6
## $ Age          <int> 34, 34, 34, 4, 49, 9, 8, 45, 45, 45, 66, 58,
54, ...
## $ Gender        <fct> male, male, male, male, female, male, male,
femal...
## $ SleepTrouble  <fct> Yes, Yes, Yes, NA, Yes, NA, NA, No, No, No, No,
N...
## $ BMI           <dbl> 32.22, 32.22, 32.22, 15.30, 30.57, 16.82,
20.64, ...
## $ HHIncome      <fct> 25000-34999, 25000-34999, 25000-34999, 20000-
2499...
## $ PhysActive    <fct> No, No, No, NA, No, NA, NA, Yes, Yes, Yes, Yes,
Y...

# Convert factors to numeric - the packages just seem to work better
that way
people2$Gender <- as.numeric(people2$Gender)
people2$SleepTrouble <- as.numeric(people2$SleepTrouble)
people2$HHIncome <- as.numeric(people2$HHIncome)
people2$PhysActive <- as.numeric(people2$PhysActive)

people2 <- na.omit(people2)

glimpse(people2)

## Observations: 7,037
## Variables: 6
## $ Age          <int> 34, 34, 34, 49, 45, 45, 45, 66, 58, 54, 58, 50,
3...
## $ Gender        <dbl> 2, 2, 2, 1, 1, 1, 1, 2, 2, 2, 1, 2, 2, 2, 1, 1,
2...
## $ SleepTrouble  <dbl> 2, 2, 2, 2, 1, 1, 1, 1, 1, 2, 1, 1, 1, 2, 1, 1,
2...
## $ BMI           <dbl> 32.22, 32.22, 32.22, 30.57, 27.24, 27.24,
27.24, ...
## $ HHIncome      <dbl> 6, 6, 6, 7, 11, 11, 11, 6, 12, 10, 11, 4, 6, 4,
1...
## $ PhysActive    <dbl> 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 1, 1, 2, 2,
2...

```

Problem 2

Apply the k-nearest neighbor procedure to predict SleepTrouble from the other covariates, as we did for SleepTrouble. Use k = 1, 3, 5, and 20.

```
#Apply k-nearest neighbor approach to predict SleepTrouble for k = 1, 3, 5, 20
```

```
# Let's try different values of k to see how that affects performance. This is taking different numbers of nearest neighbors
```

```
knn.1 <- knn(train = people2, test = people2, cl =  
as.numeric(people2$SleepTrouble), k = 1)  
knn.3 <- knn(train = people2, test = people2, cl =  
people2$SleepTrouble, k = 3)  
knn.5 <- knn(train = people2, test = people2, cl =  
people2$SleepTrouble, k = 5)  
knn.20 <- knn(train = people2, test = people2, cl =  
people2$SleepTrouble, k = 20)
```

```
#knn.1
```

Problem 3

Now let's see how well these classifiers work overall

```
# How well do these classifiers (k = 1, 3, 5, 20) work? Calculate the percent predicted correctly
```

```
100*sum(people2$SleepTrouble == knn.1)/length(knn.1)
```

```
## [1] 100
```

```
100*sum(people2$SleepTrouble == knn.3)/length(knn.3)
```

```
## [1] 91.99943
```

```
100*sum(people2$SleepTrouble == knn.5)/length(knn.5)
```

```
## [1] 88.71678
```

```
100*sum(people2$SleepTrouble == knn.20)/length(knn.20)
```

```
## [1] 78.66989
```

Similar to our in class exercise with diabetes, we see that as k increases, the prediction worsens, but this is expected. Prediction does seem to be poorer for SleepTrouble compared to Diabetes.

Problem 4

What about success overall?

```
# Another way to look at success rate against increasing k
```

```
table(knn.1, people2$SleepTrouble)
```

```
##
## knn.1      1      2
##      1 5239      0
##      2      0 1798

table(knn.3, people2$SleepTrouble)

##
## knn.3      1      2
##      1 5063  387
##      2  176 1411

table(knn.5, people2$SleepTrouble)

##
## knn.5      1      2
##      1 5030  585
##      2  209 1213

table(knn.20, people2$SleepTrouble)

##
## knn.20     1      2
##      1 5094 1356
##      2  145  442
```

This confirms what we saw earlier. $k=1$ perfectly predicts SleepTrouble and prediction worsens as k increases.

Github repository: <https://github.com/iyang5/Homework-8.git>